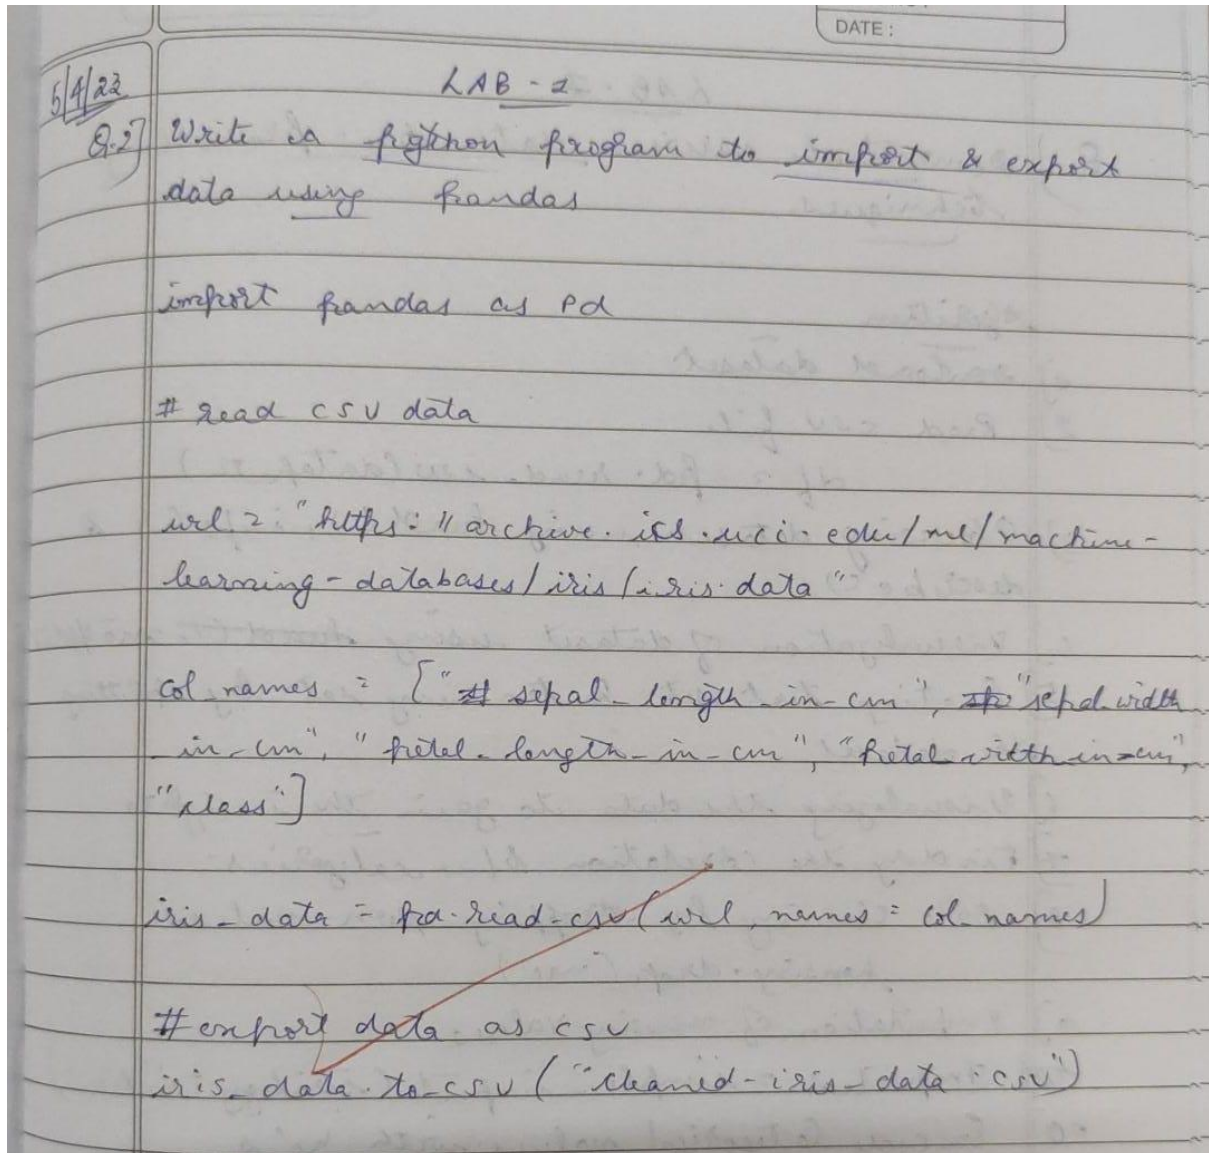


1) Write a python program to import and export data using Pandas library functions.

Importing data

Algorithm(Observation book)



Code

```
import pandas as pd
df=pd.read_csv("/content/austinHousingData.csv")
df.head(5)
```

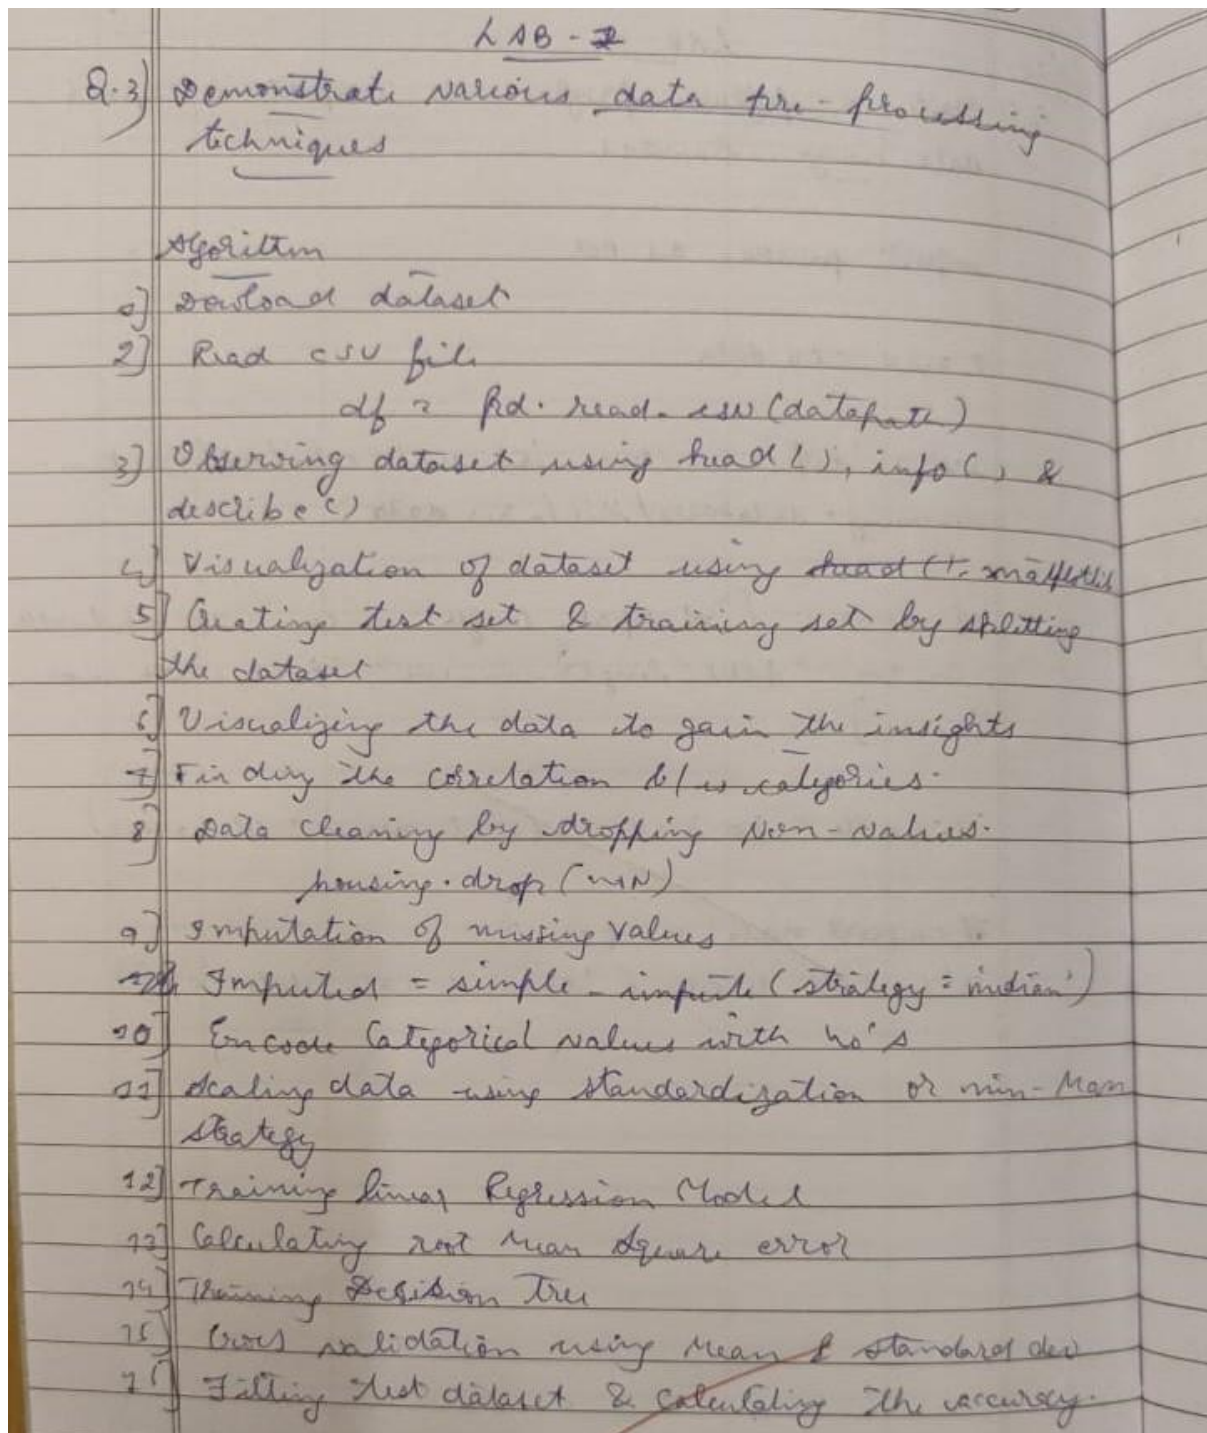
Output

	zpid	city	streetAddress	zipcode	description	latitude	longitude	propertyTaxRate	garageSpaces	hasAssociation	...
0	111373431	pflugerville	14424 Lake Victor Dr	78660	14424 Lake Victor Dr, Pflugerville, TX 78660 i...	30.430632	-97.663078	1.98	2	True	...
1	120900430	pflugerville	1104 Strickling Dr	78660	Absolutely GORGEOUS 4 Bedroom home with 2 full...	30.432673	-97.661697	1.98	2	True	...
2	2084491383	pflugerville	1408 Fort Dessau Rd	78660	Under construction - estimated completion in A...	30.409748	-97.639771	1.98	0	True	...
3	120901374	pflugerville	1025 Strickling Dr	78660	Absolutely darling one story home in charming ...	30.432112	-97.661659	1.98	2	True	...
4	60134862	pflugerville	15005 Donna Jane Loop	78660	Brimming with appeal & warm livability! Sleek ...	30.437368	-97.656860	1.98	0	True	...

5 rows × 47 columns

2) Demonstrate various data pre-processing techniques for a given dataset.

Algorithm(Observation book)



Code

```
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
```

```
# Define the column names
col_names = ["sepal_length_in_cm",
             "sepal_width_in_cm",
             "petal_length_in_cm",
             "petal_width_in_cm",
             "class"]

# Read data from URL
iris_data = pd.read_csv(url, names=col_names)
iris_data.head(5)
iris_data.to_csv("/content/exported_irisData.csv")
```

Output:

	sepal_length_in_cm	sepal_width_in_cm	petal_length_in_cm	petal_width_in_cm	class
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

```
%matplotlib inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn
df1=pd.read_csv("/content/Data.csv")
df1.head(5)
```

	Country	Age	Salary	Purchased
0	France	44.0	72000.0	No
1	Spain	27.0	48000.0	Yes
2	Germany	30.0	54000.0	No
3	Spain	38.0	61000.0	No
4	Germany	40.0	NaN	Yes

```
#Identifying and handling the missing values
df1.isnull().sum()
```

```
Country      0
Age          1
Salary       1
Purchased    0
dtype: int64
```

3) Use an appropriate dataset for building the decision tree (ID3) and apply this knowledge to classify a new sample.

Algorithm (Observation book)

Lab 2
12/4/24

Use an appropriate data set for building decision tree (ID3) & apply this knowledge to classify a new sample.

Algorithm for decision tree

```

def ID3(D, A):
    if D is pure or A is empty:
        return a leaf node with majority class in D
    else:
        A_best = argmax(Information Gain(D, A))
        root = Node(A_best)
        for v in values(A_best):
            D_v = subset(D, A_best, v)
            child = ID3(D_v, A - {A_best})
            root.add_child(v, child)
        return root
    
```

Entropy = $-P_i \log_2(P_i)$

P_i -> fraction of sample within a particular node

$I(A, D) = H(S) - \sum_{s_i} \frac{|S_v|}{|S|} * H(S_v)$

S -> Total instances
 S_v -> no. of instances for which attribute D has value v .

Code

```
# Importing the required libraries
import pandas as pd
import numpy as np
import math
data = pd.read_csv('/content/PlayTennis.csv')

def highlight(cell_value):
    '''
    Highlight yes / no values in the dataframe
    '''
    color_1 = 'background-color: pink;'
    color_2 = 'background-color: lightgreen;'

    if cell_value == 'no':
        return color_1
    elif cell_value == 'yes':
        return color_2

data.style.applymap(highlight)\
    .set_properties(subset=data.columns, **{'width': '100px'})\
    .set_table_styles([{'selector': 'th', 'props': [('background-color', 'lightgray'), ('border', '1px solid gray'), ('font-weight', 'bold')]}],
    {'selector': 'tr:hover', 'props': [('background-color', 'white'), ('border', '1.5px solid black')]}])
```

	outlook	temp	humidity	windy	play
0	sunny	hot	high	False	no
1	sunny	hot	high	True	no
2	overcast	hot	high	False	yes
3	rainy	mild	high	False	yes
4	rainy	cool	normal	False	yes
5	rainy	cool	normal	True	no
6	overcast	cool	normal	True	yes
7	sunny	mild	high	False	no
8	sunny	cool	normal	False	yes
9	rainy	mild	normal	False	yes
10	sunny	mild	normal	True	yes
11	overcast	mild	high	True	yes
12	overcast	hot	normal	False	yes
13	rainy	mild	high	True	no


```

def find_entropy(data):
    """
    Returns the entropy of the class or features
    formula:  $-\sum P(X) \log P(X)$ 
    """
    entropy = 0
    for i in range(data.nunique()):
        x = data.value_counts()[i]/data.shape[0]
        entropy += (- x * math.log(x,2))
    return round(entropy,3)

def information_gain(data, data_):
    """
    Returns the information gain of the features
    """
    info = 0
    for i in range(data_.nunique()):
        df = data[data_ == data_.unique()[i]]
        w_avg = df.shape[0]/data.shape[0]
        entropy = find_entropy(df.play)
        x = w_avg * entropy
        info += x
    ig = find_entropy(data.play) - info
    return round(ig, 3)

def entropy_and_infogain(datax, feature):
    """
    Grouping features with the same class and computing their
    entropy and information gain for splitting
    """
    for i in range(data[feature].nunique()):
        df = datax[datax[feature]==data[feature].unique()[i]]
        if df.shape[0] < 1:
            continue

        display(df[[feature, 'play']].style.applymap(highlight)\
                .set_properties(subset=[feature, 'play'], **{'width':
'80px'})\
                .set_table_styles([{'selector': 'th', 'props':
[('background-color', 'lightgray'),
('border', '1px solid gray'),
('font-weight', 'bold')]}],
                {'selector': 'td', 'props':
[('border', '1px solid gray')]}],
                {'selector': 'tr: hover', 'props':
[('background-color', 'white'),

```

```
(
'border', '1.5px solid black']]]))

    print(f'Entropy of {feature} - {data[feature].unique()[i]} =
{find_entropy(df.play)}')
    print(f'Information Gain for {feature} = {information_gain(datax,
datax[feature])}')

```

```
#Computing entropy for the entire dataset
print(f'Entropy of the entire dataset: {find_entropy(data.play)}')
```

➞ Entropy of the entire dataset: 0.94

```
#Calculate the Information Gain for each feature.
```

```
#Outlook
```

```
entropy_and_infogain(data, 'outlook')
```

➞

	outlook	play
0	sunny	no
1	sunny	no
7	sunny	no
8	sunny	yes
10	sunny	yes

Entropy of outlook - sunny = 0.971

	outlook	play
2	overcast	yes
6	overcast	yes
11	overcast	yes
12	overcast	yes

Entropy of outlook - overcast = 0.0

	outlook	play
3	rainy	yes
4	rainy	yes
5	rainy	no
9	rainy	yes
13	rainy	no

Entropy of outlook - rainy = 0.971

Information Gain for outlook = 0.246


```
#Temp
```

```
entropy_and_infogain(data, 'temp')
```



	temp	play
0	hot	no
1	hot	no
2	hot	yes
12	hot	yes

Entropy of temp - hot = 1.0

	temp	play
3	mild	yes
7	mild	no
9	mild	yes
10	mild	yes
11	mild	yes
13	mild	no

Entropy of temp - mild = 0.918

	temp	play
4	cool	yes
5	cool	no
6	cool	yes
8	cool	yes

Entropy of temp - cool = 0.811

Information Gain for temp = 0.029

```
#Humidity
```

```
entropy_and_infogain(data, 'humidity')
```



	humidity	play
0	high	no
1	high	no
2	high	yes
3	high	yes
7	high	no
11	high	yes
13	high	no

Entropy of humidity - high = 0.985

	humidity	play
4	normal	yes
5	normal	no
6	normal	yes
8	normal	yes
9	normal	yes
10	normal	yes
12	normal	yes

Entropy of humidity - normal = 0.592

Information Gain for humidity = 0.151

```
#Windy
entropy_and_infogain(data, 'windy')
```



	windy	play
0	False	no
2	False	yes
3	False	yes
4	False	yes
7	False	no
8	False	yes
9	False	yes
12	False	yes

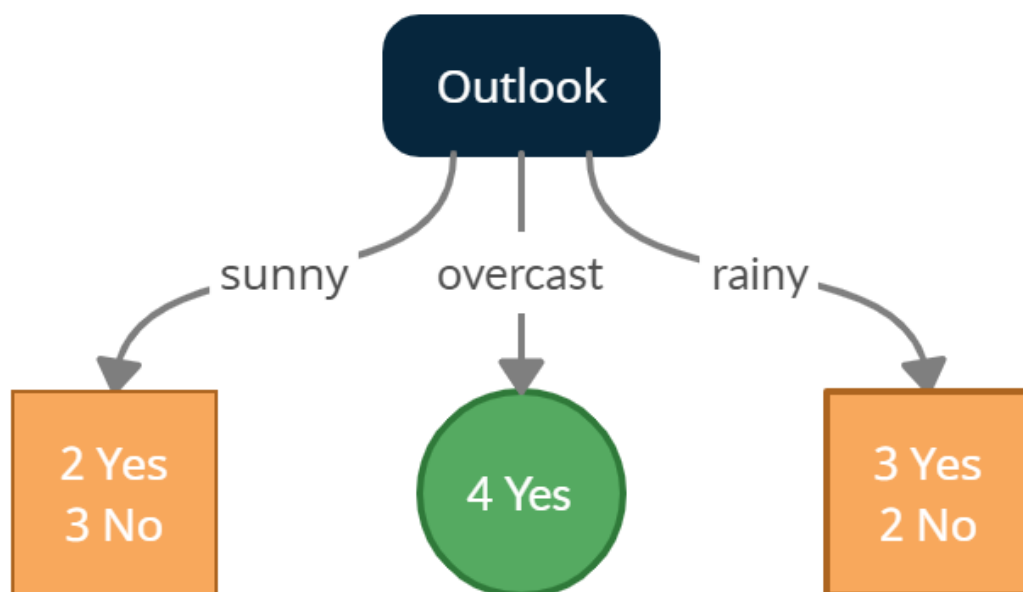
Entropy of windy - False = 0.811

	windy	play
1	True	no
5	True	no
6	True	yes
10	True	yes
11	True	yes
13	True	no

Entropy of windy - True = 1.0

Information Gain for windy = 0.048

#Make a decision tree node using the feature with the maximum Information Gain.



```

sunny = data[data['outlook'] == 'sunny']
sunny.style.applymap(highlight)\
    .set_properties(subset=data.columns, **{'width': '100px'})\
    .set_table_styles([{'selector': 'th', 'props': [('background-color', 'lightgray'), ('border', '1px solid gray'), ('font-weight', 'bold')]}],
    {'selector': 'tr:hover', 'props': [('background-color', 'white'), ('border', '1.5px solid black')]}])

```



	outlook	temp	humidity	windy	play
0	sunny	hot	high	False	no
1	sunny	hot	high	True	no
7	sunny	mild	high	False	no
8	sunny	cool	normal	False	yes
10	sunny	mild	normal	True	yes

```

print(f'Entropy of the Sunny dataset: {find_entropy(sunny.play)}')
Entropy of the Sunny dataset: 0.971

```

#temp

```

entropy_and_infogain(sunny, 'temp')

```

	temp	play
0	hot	no
1	hot	no

Entropy of temp - hot = 0.0

	temp	play
7	mild	no
10	mild	yes

Entropy of temp - mild = 1.0

	temp	play
8	cool	yes

Entropy of temp - cool = 0.0

Information Gain for temp = 0.571

```
#Humidity
```

```
entropy_and_infogain(sunny, 'humidity')
```

	humidity	play
0	high	no
1	high	no
7	high	no

Entropy of humidity - high = 0.0

	humidity	play
8	normal	yes
10	normal	yes

Entropy of humidity - normal = 0.0

Information Gain for humidity = 0.971

```
#Windy
```

```
entropy_and_infogain(sunny, 'windy')
```

	windy	play
0	False	no
7	False	no
8	False	yes

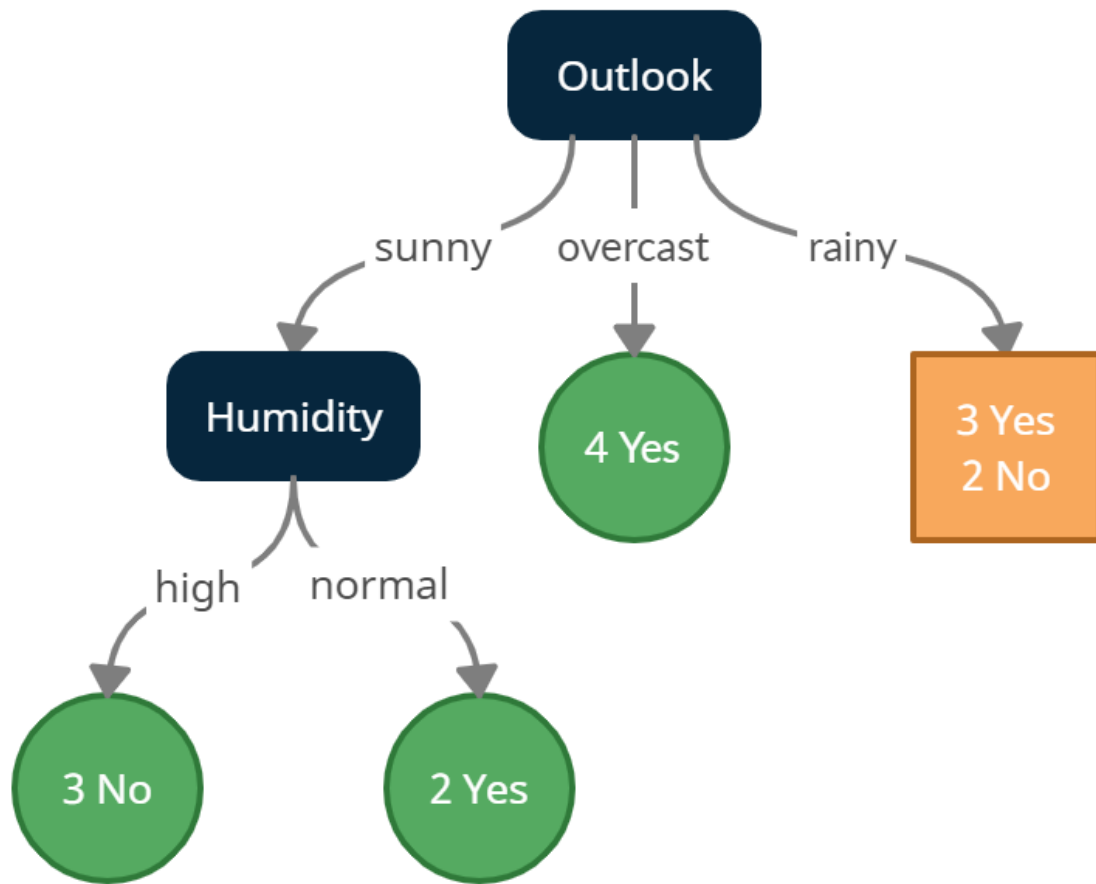
Entropy of windy - False = 0.918

	windy	play
1	True	no
10	True	yes

Entropy of windy - True = 1.0

Information Gain for windy = 0.02

#Making a decision tree node using the feature which has the maximum Information Gain



```
#Outlook - Rainy
rainy = data[data['outlook'] == 'rainy']
rainy.style.applymap(highlight)\
    .set_properties(subset=data.columns, **{'width': '100px'})\
    .set_table_styles([{'selector': 'th', 'props': [('background-color', 'lightgray'), ('border', '1px solid gray'), ('font-weight', 'bold')]}],
    {'selector': 'tr:hover', 'props': [('background-color', 'white'), ('border', '1.5px solid black')]}])
```



	outlook	temp	humidity	windy	play
3	rainy	mild	high	False	yes
4	rainy	cool	normal	False	yes
5	rainy	cool	normal	True	no
9	rainy	mild	normal	False	yes
13	rainy	mild	high	True	no

```
print(f'Entropy of the Rainy dataset: {find_entropy(rainy.play)}')
```

Entropy of the Rainy dataset: 0.971

```
#temp
```

```
entropy_and_infogain(rainy, 'temp')
```

	temp	play
3	mild	yes
9	mild	yes
13	mild	no

Entropy of temp - mild = 0.918

	temp	play
4	cool	yes
5	cool	no

Entropy of temp - cool = 1.0

Information Gain for temp = 0.02

```
#Humidity
```

```
entropy_and_infogain(rainy, 'humidity')
```

	humidity	play
3	high	yes
13	high	no

Entropy of humidity - high = 1.0

	humidity	play
4	normal	yes
5	normal	no
9	normal	yes

Entropy of humidity - normal = 0.918

Information Gain for humidity = 0.02

```
#Windy
```

```
entropy_and_infogain(rainy, 'windy')
```



	windy	play
3	False	yes
4	False	yes
9	False	yes


Entropy of windy - False = 0.0

	windy	play
5	True	no
13	True	no

Entropy of windy - True = 0.0

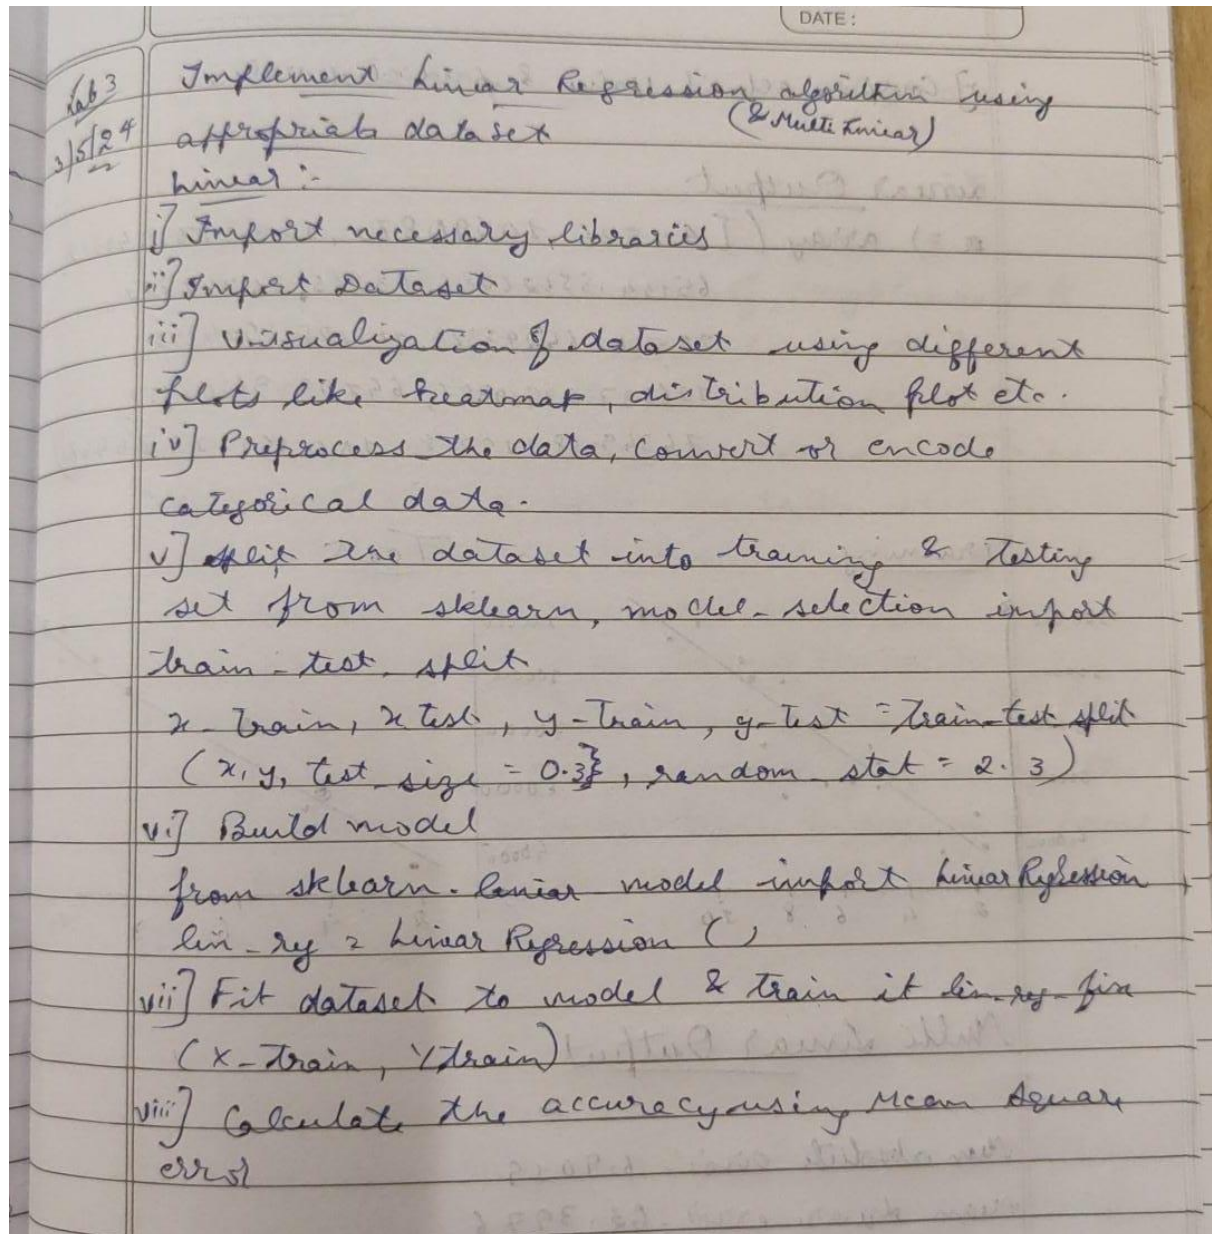
Information Gain for windy = 0.971

#Making a decision tree node using the feature which has the maximum Information Gain.



4) Implement Linear and Multi-Linear Regression algorithm using appropriate dataset

Algorithm



Multi Linear

i - iii) same

iv) Encode Categorical Data

if `ct = ColumnTransformer(transformers=[('encode', OneHotEncoder[3])], remainder='passthrough')`

v) Split dataset into training & testing set

vi) We can see multiple independent variables.

vii) Create Regression model

`Regressor = LinearRegression()`

viii) Fit Train set

ix) Test model using test set

Code:

```
#Importing the libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# import warnings
import warnings
warnings.filterwarnings("ignore")

# We will use some methods from the sklearn module
from sklearn import linear_model
from sklearn.linear_model import LinearRegression
from sklearn import metrics
from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.model_selection import train_test_split, cross_val_score
```

```
# Reading the Dataset
df = pd.read_csv("data.csv")
```

```
df.head()
```

	Car	Model	Volume	Weight	CO2
0	Toyoty	Aygo	1000	790	99
1	Mitsubishi	Space Star	1200	1160	95
2	Skoda	Citigo	1000	929	95
3	Fiat	500	900	865	90
4	Mini	Cooper	1500	1140	105

```
df.shape
```

```
(36, 5)
```

```
df.corr(numeric_only=True)
```


	Volume	Weight	CO2
Volume	1.000000	0.753537	0.592082
Weight	0.753537	1.000000	0.552150
CO2	0.592082	0.552150	1.000000

```
In [ ]: print(df.describe())
```


	Volume	Weight	CO2
count	36.000000	36.000000	36.000000
mean	1611.111111	1292.277778	102.027778
std	388.975047	242.123889	7.454571
min	900.000000	790.000000	90.000000
25%	1475.000000	1117.250000	97.750000
50%	1600.000000	1329.000000	99.000000
75%	2000.000000	1418.250000	105.000000
max	2500.000000	1746.000000	120.000000

```
In [ ]: #Setting the value for X and Y
x = df[['Weight', 'Volume']]
y = df['CO2']
```


```
In [ ]: fig, axs = plt.subplots(2, figsize = (5,5))
plt1 = sns.boxplot(df['Weight'], ax = axs[0])
plt2 = sns.boxplot(df['Volume'], ax = axs[1])
plt.tight_layout()
```



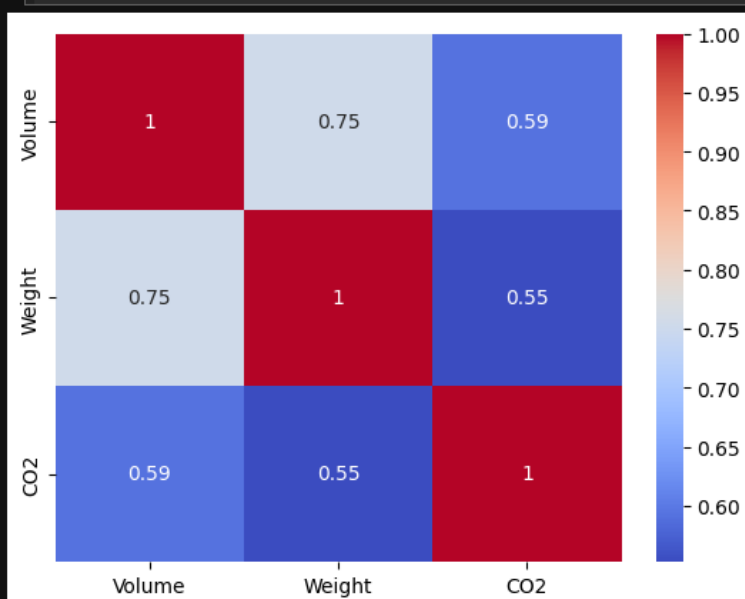
```
In [ ]: sns.distplot(df['CO2']);
```



```
In [ ]: sns.pairplot(df, x_vars=['Weight', 'Volume'], y_vars='CO2', height=4, aspect=1, kind='scatter')
plt.show()
```



```
In [ ]: # Create the correlation matrix and represent it as a heatmap.
sns.heatmap(df.corr(numeric_only=True), annot = True, cmap = 'coolwarm')
plt.show()
```



```
In [ ]: X_train,X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 100)
```

```
In [ ]: y_train.shape
```

```
Out[ ]: (25,)
```

```
In [ ]: y_test.shape
```

```
Out[ ]: (11,)
```

```
In [ ]: reg_model = linear_model.LinearRegression()
```

```
In [ ]: #Fitting the Multiple Linear Regression model  
reg_model = LinearRegression().fit(X_train, y_train)
```

```
In [ ]: #Printing the model coefficients  
print('Intercept: ',reg_model.intercept_)  
# pair the feature names with the coefficients  
list(zip(X, reg_model.coef_))
```

```
Intercept: 74.33882836589245
```

```
Out[ ]: [('Weight', 0.0171800645996374), ('Volume', 0.0025046399866402976)]
```

```
In [ ]: #Predicting the Test and Train set result  
y_pred= reg_model.predict(X_test)  
x_pred= reg_model.predict(X_train)
```

```
In [ ]: print("Prediction for test set: {}".format(y_pred))
```

```
Prediction for test set: [ 90.41571939 102.16323413  99.56363213 104.56661845 101.54657652  
95.94770019 108.64011848 102.22654214  92.80374837  97.27327129  
97.57074463]
```

```
In [ ]: #Actual value and the predicted value  
reg_model_diff = pd.DataFrame({'Actual value': y_test, 'Predicted value': y_pred})  
reg_model_diff
```

```
Out[ ]: 

|    | Actual value | Predicted value |
|----|--------------|-----------------|
| 0  | 99           | 90.415719       |
| 19 | 105          | 102.163234      |
| 32 | 104          | 99.563632       |
| 35 | 120          | 104.566618      |
| 7  | 92           | 101.546577      |
| 12 | 99           | 95.947700       |
| 29 | 114          | 108.640118      |
| 33 | 108          | 102.226542      |
| 5  | 105          | 92.803748       |
| 1  | 95           | 97.273271       |
| 18 | 104          | 97.570745       |


```

```
In [ ]: mae = metrics.mean_absolute_error(y_test, y_pred)
mse = metrics.mean_squared_error(y_test, y_pred)
r2 = np.sqrt(metrics.mean_squared_error(y_test, y_pred))

print('Mean Absolute Error:', mae)
print('Mean Square Error:', mse)
print('Root Mean Square Error:', r2)
```

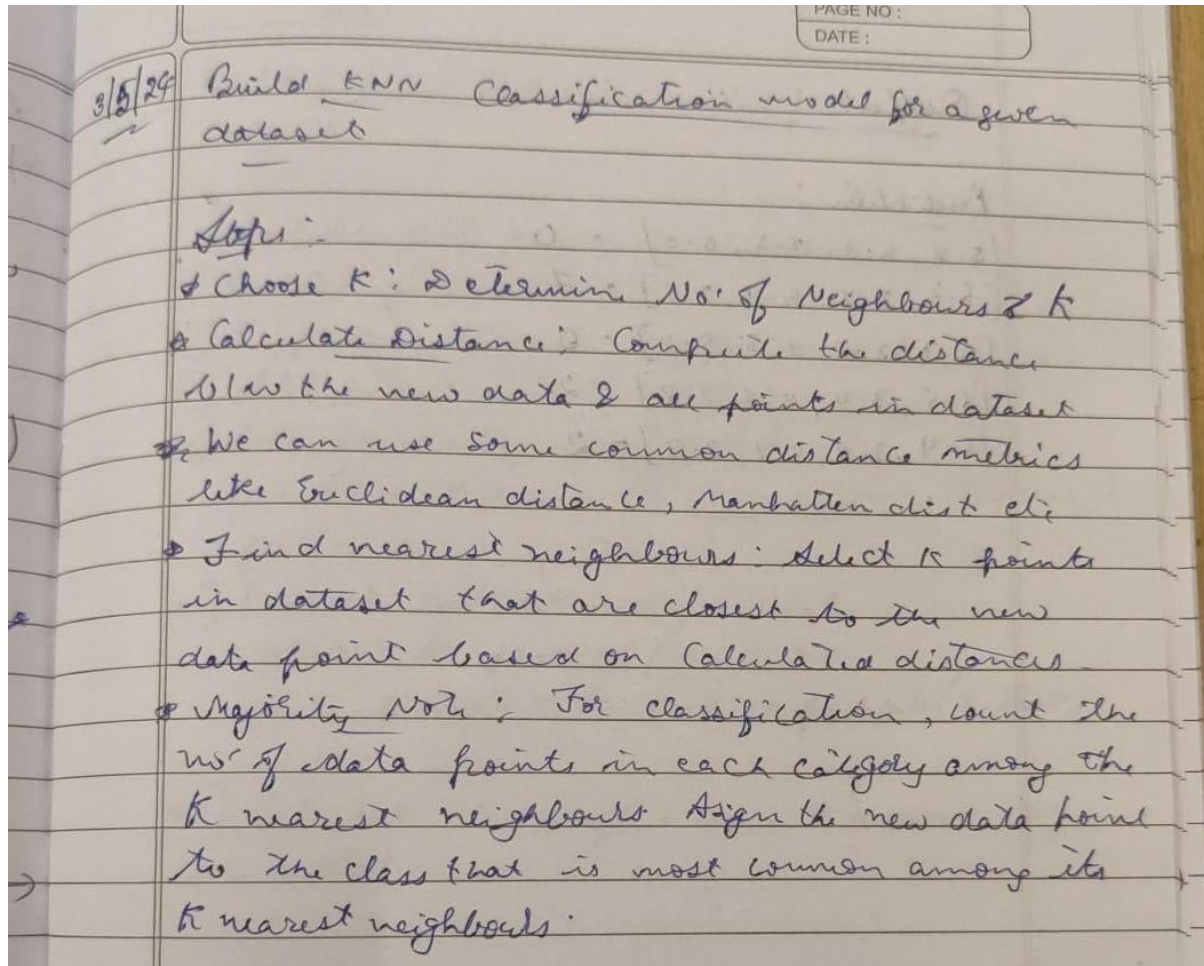
```
Mean Absolute Error: 6.901980901636316
Mean Square Error: 63.39765310998794
Root Mean Square Error: 7.96226432053018
```

Results

```
Mean Absolute Error: 6.901980901636316
Mean Square Error: 63.39765310998794
Root Mean Square Error: 7.96226432053018
```


5) Build KNN Classification model for a given dataset.

Algorithm



Code:

```
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix
from sklearn import datasets
iris = datasets.load_iris()

x = iris.data
y = iris.target

print('sepal-length', 'sepal-width', 'petal-length', 'petal-width')
print(x)
print('class: 0 - Iris-Setosa, 1 - Iris-Versicolour, 2 - Iris-Virginica')
print(y)
```

```

In [5]: x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.3)

#To Training the model and Nearest nighbors K=5
classifier = KNeighborsClassifier(n_neighbors=5)
classifier.fit(x_train, y_train)

#to make predictions on our test data
y_pred=classifier.predict(x_test)

print('Prediction -')

for i,test in enumerate(x_test) :
    print(f'{test} - {y_pred[i]}')

# print('Confusion Matrix')
# print(confusion_matrix(y_test,y_pred))
# print('Accuracy Metrics')
# print(classification_report(y_test,y_pred))

```

Results:

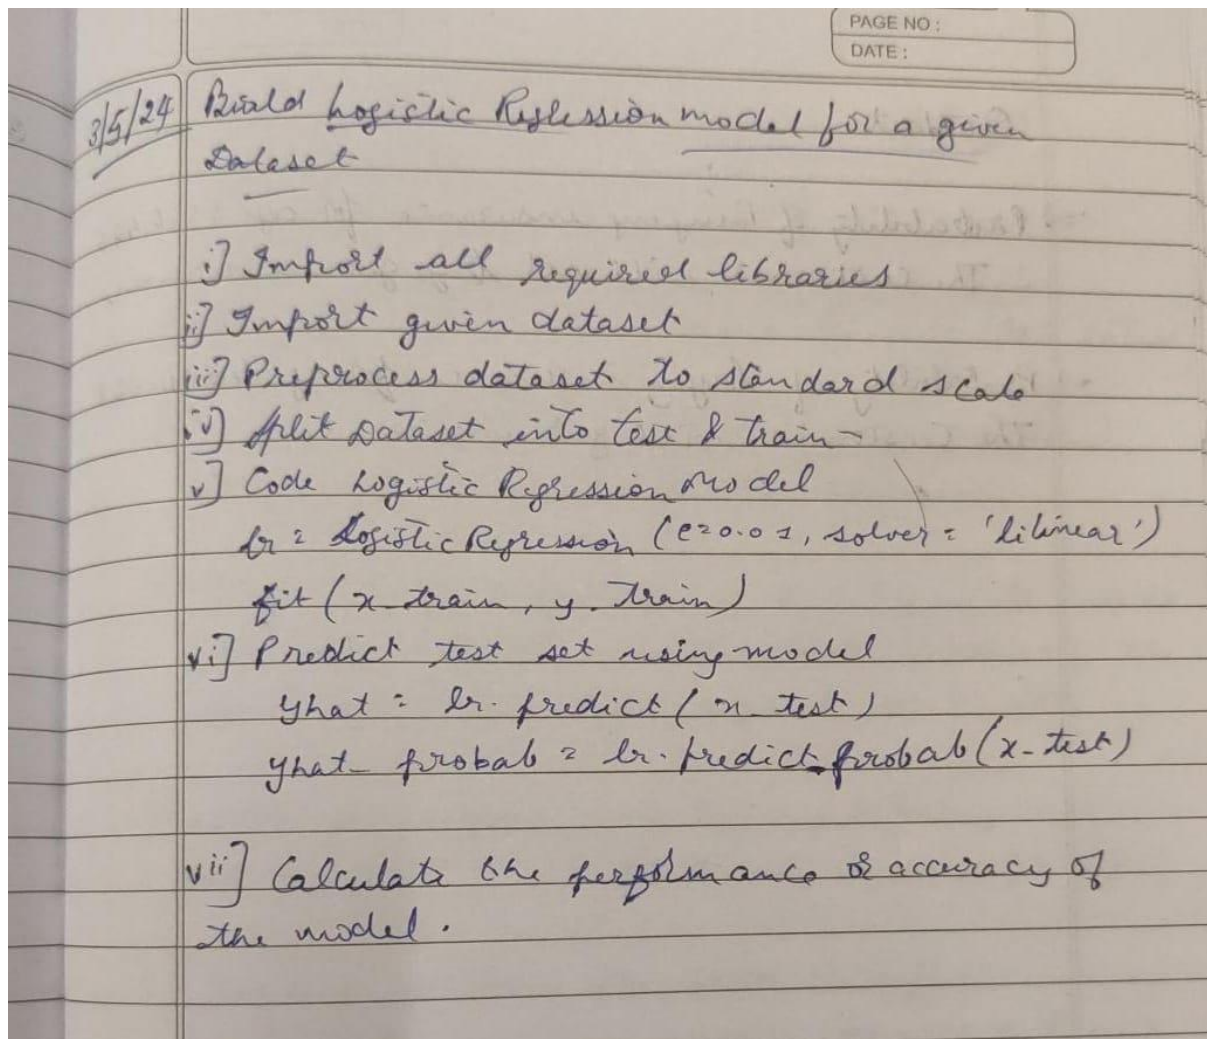
```

Prediction -
[5.2 4.1 1.5 0.1] - 0
[5.5 2.3 4.  1.3] - 1
[6.7 3.1 4.7 1.5] - 1
[7.  3.2 4.7 1.4] - 1
[6.2 2.8 4.8 1.8] - 2
[5.7 2.8 4.5 1.3] - 1
[6.  3.4 4.5 1.6] - 1
[5.1 3.8 1.6 0.2] - 0
[5.5 2.5 4.  1.3] - 1
[4.8 3.1 1.6 0.2] - 0
[6.1 3.  4.9 1.8] - 2
[4.7 3.2 1.6 0.2] - 0
[5.6 2.9 3.6 1.3] - 1
[5.4 3.9 1.3 0.4] - 0
[5.  3.2 1.2 0.2] - 0
[6.1 2.9 4.7 1.4] - 1
[5.  3.4 1.5 0.2] - 0
[7.7 2.8 6.7 2.  ] - 2
[4.6 3.2 1.4 0.2] - 0
[5.7 2.9 4.2 1.3] - 1
[4.6 3.6 1.  0.2] - 0
[6.8 2.8 4.8 1.4] - 1
[6.8 3.2 5.9 2.3] - 2

```

6) Build Logistic Regression Model for a given dataset

Algorithm:




Code:

```
In [1]: import pandas as pd
        from matplotlib import pyplot as plt
        %matplotlib inline
```

```
In [4]: df = pd.read_csv("insurance_data.csv")
        df.head()
        plt.scatter(df.age,df.bought_insurance,marker = '+',color = 'red')
```

```
Out[4]: <matplotlib.collections.PathCollection at 0x7e3b7e4c9cf0>
```



```
In [31]: from sklearn.model_selection import train_test_split
        X_train, X_test, y_train, y_test = train_test_split(df[['age']],df.bought_insurance,train_size=0.7)
        print(X_test)
```

```
age
12  27
13  29
4   46
23  45
3   52
8   62
1   25
11  28
25  54
```

```
In [32]: from sklearn.linear_model import LogisticRegression
        model = LogisticRegression()
        model.fit(X_train, y_train)
        print(X_test)
        y_predicted = model.predict(X_test)
        probab = model.predict_proba(X_test)
        score = model.score(X_test,y_test)
        print(y_predicted)
        print("\nprobability: ")
        print(probab)
        print("\naccuracy: ")
        print(score)
```

```
age
12  27
13  29
4   46
23  45
3   52
8   62
1   25
11  28
25  54
[0 0 1 1 1 1 0 0 1]
```

```
probability:
[[0.75147045 0.24852955]
 [0.69990447 0.30009553]
 [0.20424226 0.79575774]
 [0.2261509  0.7738491 ]
 [0.10537592 0.89462408]
 [0.03115876 0.96884124]
 [0.79674889 0.20325111]
 [0.7264443  0.2735557 ]
 [0.08328741 0.91671259]]
```

```
accuracy:
0.8888888888888888
```

```
In [33]: print(X_test)
```

```
age
12  27
13  29
4   46
23  45
3   52
8   62
1   25
11  28
25  54
```

```
In [38]: import math
def sigmoid(x):
    return 1 / (1 + math.exp(-x))
```

```
In [39]: def prediction_function(age):
z = 0.042 * age - 1.53 # 0.04150133 ~ 0.042 and -1.52726963 ~ -1.53
y = sigmoid(z)
return y
```

```
In [40]: def insure(probability):
    if (probability > 0.5):
        print("The customer will get insurance.")
    else:
        print("The customer will not get insurance.")

age = 35
probability = prediction_function(age)
print("Probability of buying insurance for age 35:", probability)
insure(probability)
```

```
Probability of buying insurance for age 35: 0.4850044983805899
The customer will not get insurance.
```

```
In [41]: age = 43
probability = prediction_function(age)
print("Probability of buying insurance for age 43:", probability)
insure(probability)
```

```
Probability of buying insurance for age 43: 0.568565299077705
The customer will get insurance.
```

Results:

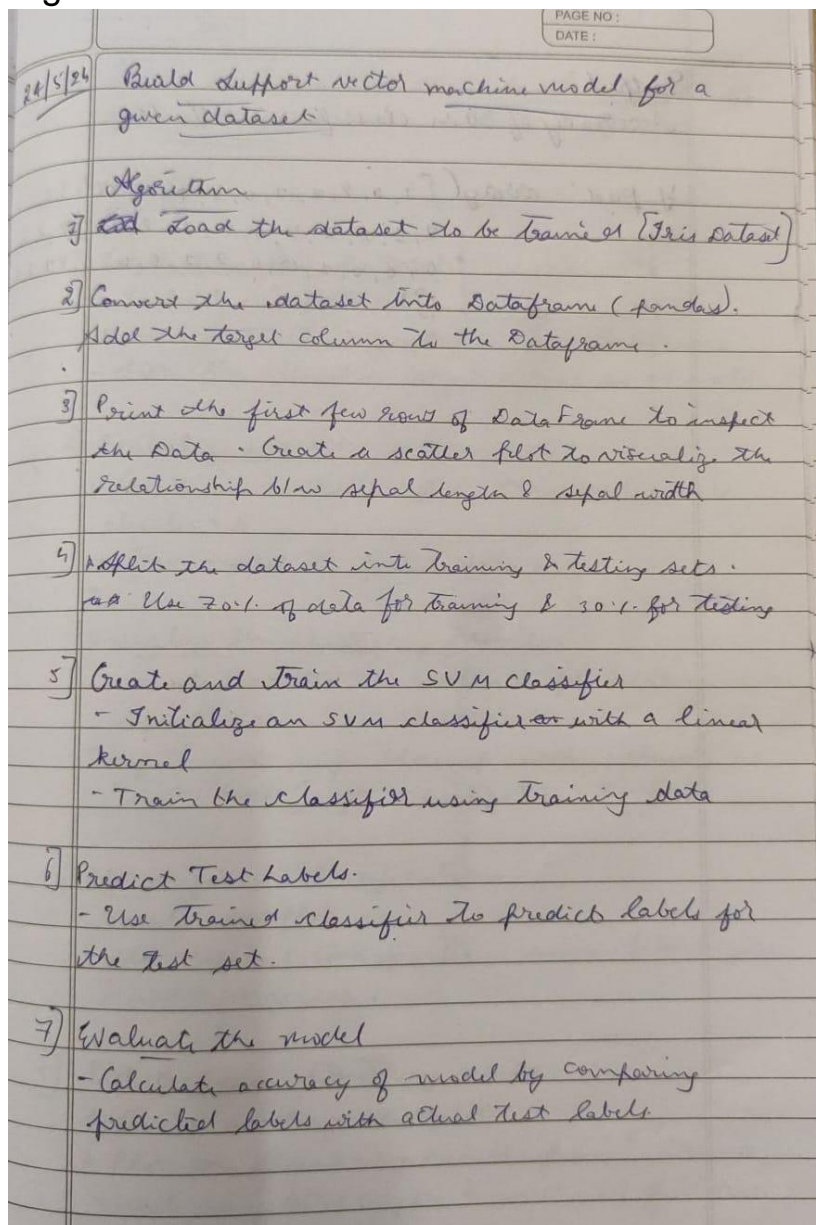
```
Probability of buying insurance for age 35: 0.4850044983805899
The customer will not get insurance.

In [41]: age = 43
          probability = prediction_function(age)
          print("Probability of buying insurance for age 43:", probability)
          insure(probability)

Probability of buying insurance for age 43: 0.568565299077705
The customer will get insurance.
```

7) Build Support vector machine model for a given dataset

Algorithm:



Code:

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
```

```
In [2]: # Load the Iris dataset
iris = load_iris()
```

```
In [3]: # Convert the dataset into a pandas DataFrame
iris_df = pd.DataFrame(data=iris.data, columns=iris.feature_names)
iris_df['target'] = iris.target
```


```
In [4]: # Display the first few rows of the DataFrame
print(iris_df.head())
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	\
0	5.1	3.5	1.4	0.2	
1	4.9	3.0	1.4	0.2	
2	4.7	3.2	1.3	0.2	
3	4.6	3.1	1.5	0.2	
4	5.0	3.6	1.4	0.2	

	target
0	0
1	0
2	0
3	0
4	0

In [5]:

```
# Plotting to visualize the data
plt.figure(figsize=(10, 6))
plt.scatter(iris_df['sepal length (cm)'], iris_df['sepal width (cm)'],
            c=iris_df['target'], cmap='viridis')
plt.xlabel('Sepal Length (cm)')
plt.ylabel('Sepal Width (cm)')
plt.title('Iris Dataset - Sepal Length vs Sepal Width')
plt.colorbar(label='Species')
plt.show()
```



In [6]:

```
# Splitting the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(iris.data, iris.target, test_size=0.3, random_state=42)
```

In [7]:

```
# Creating and training the SVM classifier
svm_classifier = SVC(kernel='linear')
svm_classifier.fit(X_train, y_train)

# Predicting the labels for the test set
y_pred = svm_classifier.predict(X_test)

# Calculating the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy of SVM Classifier:", accuracy)
```

Results:

Accuracy of SVM Classifier: 1.0

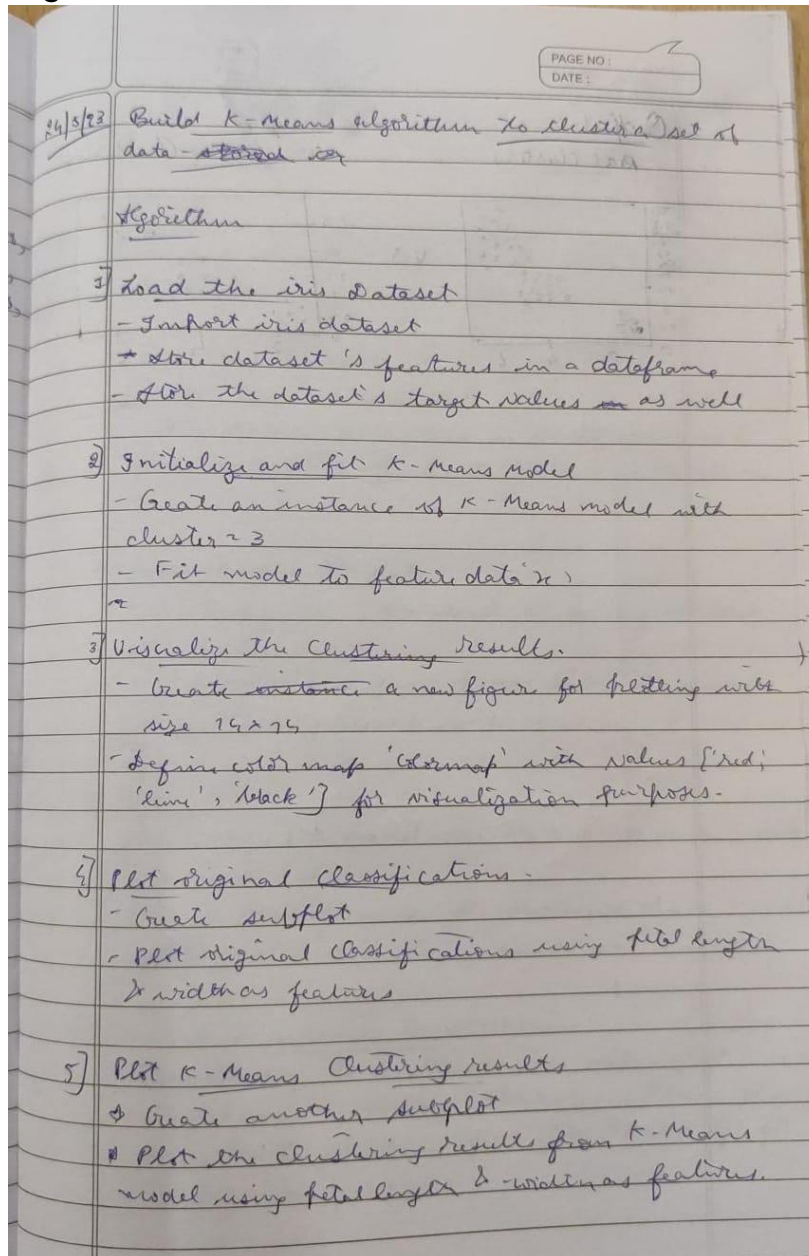
In [8]:

y_pred

Out[8]: array([1, 0, 2, 1, 1, 0, 1, 2, 1, 1, 2, 0, 0, 0, 0, 1, 2, 1, 1, 2, 0, 2,
0, 2, 2, 2, 2, 2, 0, 0, 0, 0, 1, 0, 0, 2, 1, 0, 0, 0, 2, 1, 1, 0,
0])

8) Build k-Means algorithm to cluster a set of data stored in a .CSV file.

Algorithm:



Code:

```

import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.cluster import KMeans
import pandas as pd
import numpy as np

[] # Import some data to play with
iris = datasets.load_iris()
X = pd.DataFrame(iris.data)
X.columns = ['Sepal_Length', 'Sepal_Width', 'Petal_Length', 'Petal_Width']
y = pd.DataFrame(iris.target)
y.columns = ['Targets']

[] # Build the K Means Model
model = KMeans(n_clusters=3)
model.fit(X) # model.labels_ gives cluster no for which samples belongs

/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of 'n_init' will change from 10 to 'auto' in 1.4. Set the value of 'n_init' explicitly to suppress the warning
  warnings.warn(
    KMeans
    KMeans(n_clusters=3)


[] # Visualise the clustering results
plt.figure(figsize=(14,14))
colormap = np.array(['red', 'lime', 'black'])

<Figure size 1400x1400 with 0 Axes>

[] # Plot the Original Classifications using Petal features
plt.subplot(2, 2, 1)
plt.scatter(X.Petal_Length, X.Petal_Width, c=colormap[y.Targets], s=40)
plt.title('Real Clusters')
plt.xlabel('Petal Length')
plt.ylabel('Petal Width')

```

Text(0, 0.5, 'Petal Width')




```

# Plot the Models Classifications
plt.subplot(2, 2, 2)
plt.scatter(X.Petal_Length, X.Petal_Width, c=colormap[model.labels_], s=40)
plt.title('K-Means Clustering')
plt.xlabel('Petal Length')
plt.ylabel('Petal Width')

```

Text(0, 0.5, 'Petal Width')



9) Implement Dimensionality reduction using Principle Component Analysis (PCA) method.

Algorithm

PAGE NO:
 DATE:

20/1/24 Implement Dimensionality Reduction using Principle Component Analysis method

Algorithm

- 1] Import necessary libraries for data handling, standardization, PCA & plotting
- 2] Load Iris Dataset
 - Load the iris dataset.
 - Store dataset's features in DataFrame x & store dataset's target values in DataFrame y .
- 3] Standardize the data
 - Initialize a 'StandardScaler' to standardize the feature data.
 - Fit & transform feature data x using scaler, storing result in x_scaled .
- 4] Apply PCA
 - Initialize PCA model with $n=2$ to reduce data to 2 dimensions.
 - Fit & transform the standardized data x_scaled using PCA model storing result in x_pca .
- 5] Convert PCA result to DataFrame
 - Convert the PCA transformed data x_pca into a DataFrame x_pca_df with columns ['PC1', 'PC2']
 - Add the target column from y to x_pca_df for visualization purposes.

DATE:

32/

- 6] Visualize the PCA result
 - Create a new figure for plotting
 - Define a color map array 'colormap' with values ['red', 'lime', 'black']
 - Create scatter plot of PCA-transformed data using principal components.
 - Color points based on original class labels from y .

Code:

```
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
import pandas as pd
import numpy as np

# Load the iris dataset
iris = datasets.load_iris()
X = pd.DataFrame(iris.data, columns=['Sepal_Length', 'Sepal_Width', 'Petal_Length', 'Petal_Width'])
y = pd.DataFrame(iris.target, columns=['Targets'])

# Standardize the data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Apply PCA
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)


# Convert PCA result to a DataFrame
X_pca_df = pd.DataFrame(X_pca, columns=['PCA1', 'PCA2'])

# Add the target column for visualization
X_pca_df['Targets'] = y.Targets

# Visualize the PCA result
plt.figure(figsize=(14, 7))
colormap = np.array(['red', 'lime', 'black'])

# Plot the PCA transformed data
plt.scatter(X_pca_df.PCA1, X_pca_df.PCA2, c=colormap[X_pca_df.Targets], s=40)
plt.title('PCA of Iris Dataset')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()
```

Results:



10) Build Artificial Neural Network model with back propagation on a given dataset

Algorithm

31/5/24 Build Artificial Neural Network model with back propagation

Algorithm

- 1] Initialize parameters
 - Normalize i/p feature matrix 'x'
 - Normalize the output 'y'
 - Set hyperparameters: no. of epochs, lr, no. of neurons in i/p layer, hidden layer & o/p layer
 - Initialize weights & biases for hidden layer & o/p layer with random values.
- 2] Define Activation Functions
 - Sigmoid $\sigma(x) = \frac{1}{1+e^{-x}}$
 - Derivative of Sigmoid $\sigma'(x) = \sigma(x) \cdot (1-\sigma(x))$
- 3] Training the Network

Forward Propagation

 - Compute i/p to hidden layer: $h_{inpl} = x \cdot w_{ih}$
 - Add bias to hidden layer i/p: $h_{inpl} = h_{inpl} + b_{ih}$
 - Apply Activation function: $h_{layer_act} = \sigma(h_{inpl})$
 - Compute Activation function i/p to o/p layer:
 $out_{inp} = h_{layer_act} \cdot w_{oh}$
 - Add bias to output layer input: $out_{inp} = out_{inp} + b_{oh}$
 - Apply activation fn: $o/p = \sigma(out_{inp})$

Backpropagation

- Compute error at o/p layer: $E_O = y - output$
- Compute gradient at o/p layer: $outgrad = \sigma'(output)$
- Compute delta for o/p layer: $d_{output} = E_O \cdot outgrad$
- Compute error at hidden layer: $E_H = d_{output} \cdot w_{oh}$
- Compute gradient at hidden layer: $hiddenlayer = \sigma'(h_{layer_act})$
- Compute delta for hidden layer: $d_{hiddenlayer} = E_H \cdot hiddenlayer$

Update wts & biases

- Update wts for o/p layer: $w_{oh} += \eta \cdot h_{layer_act} \cdot d_{output}$
- $b_{oh} += \eta \cdot d_{output}$
- Update wts for hidden layer: $w_{ih} += \eta \cdot x \cdot d_{hiddenlayer}$
- $b_{ih} += \eta \cdot d_{hiddenlayer}$

Code:

```
In [10]: import numpy as np
x = np.array([[2,9],[1,5],[3,6]],dtype = float)
y = np.array([[92],[86],[89]],dtype = float)

x = x/np.amax(x,axis=0)
y = y/100

In [11]: #Variable Initialization
epoch = 5000
lr = 0.1
inputlayer_neurons = 2
hiddenlayer_neurons = 3
output_neurons = 1

In [12]: # weight and bias Initialization

wh = np.random.uniform(size=(inputlayer_neurons,hiddenlayer_neurons))
bh = np.random.uniform(size=(1,hiddenlayer_neurons))
wout = np.random.uniform(size=(hiddenlayer_neurons,output_neurons))
bout = np.random.uniform(size=(1,output_neurons))

In [13]: #sigmoid function

def sigmoid(x):
    return 1/(1+np.exp(-x))

# Derivative of Sigmoid

def der_sigmoid(x):
    return x*(1-x)
```

```
In [14]: # Draws a random range of numbers uniformly of dim x*y

for i in range(epoch):

    # forward propagation
    hinp1 = np.dot(x,wh)
    hinp = hinp1 + bh
    hlayer_act = sigmoid(hinp)
    outinp1 = np.dot(hlayer_act,wout)
    outinp = outinp1 + bout
    output = sigmoid(outinp)

    # Backpropagation
    E0 = y - output
    outgrad = der_sigmoid(output)
    d_output = E0*outgrad
    EH = d_output.dot(wout.T)

In [15]: # how much hidden layer weights contributed to error

hiddengrad = der_sigmoid(hlayer_act)
d_hiddenlayer = EH*hiddengrad

#dotproduct of nextlayererror and current layer op

wout += hlayer_act.T.dot(d_output)*lr
wh += x.T.dot(d_hiddenlayer)*lr

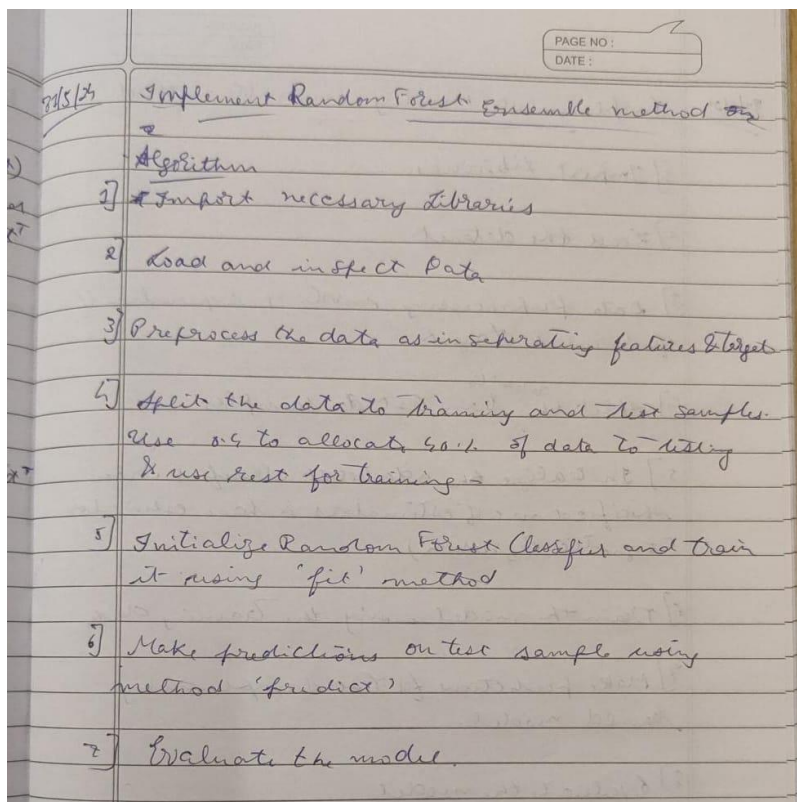
print("Input: \n" + str(x))
print("Actual output: \n" + str(y))
print("Predicted Output: \n",output)
```


Results

```
Input:
[[0.66666667 1.      ]
 [0.33333333 0.55555556]
 [1.      0.66666667]]
Actual output:
[[0.92]
 [0.86]
 [0.89]]
Predicted Output:
[[0.80056875]
 [0.79393831]
 [0.80112347]]
```

- 11) Implement Random forest ensemble method on a given dataset.

Algorithm



Code:

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn import datasets

# Load the data
iris_data = datasets.load_iris()

X = pd.DataFrame(iris_data.data, columns=['Sepal_Length', 'Sepal_Width', 'Petal_Length', 'Petal_Width'])
y = pd.DataFrame(iris_data.target, columns=['Targets'])

# Check the info of the modified data
# print(iris_data.info())

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random_state=42)

# Initialize the Random Forest classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)

# Fit the classifier to the training data
rf_classifier.fit(X_train, y_train)

# Predict on the test data
y_pred = rf_classifier.predict(X_test)

# Evaluate the classifier
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")

# Print classification report
print("Classification Report:")
print(classification_report(y_test, y_pred))

# Print confusion matrix
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

```

Results:

```

Accuracy: 0.98
Classification Report:

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	23
1	0.95	1.00	0.97	19
2	1.00	0.94	0.97	18
accuracy			0.98	60
macro avg	0.98	0.98	0.98	60
weighted avg	0.98	0.98	0.98	60

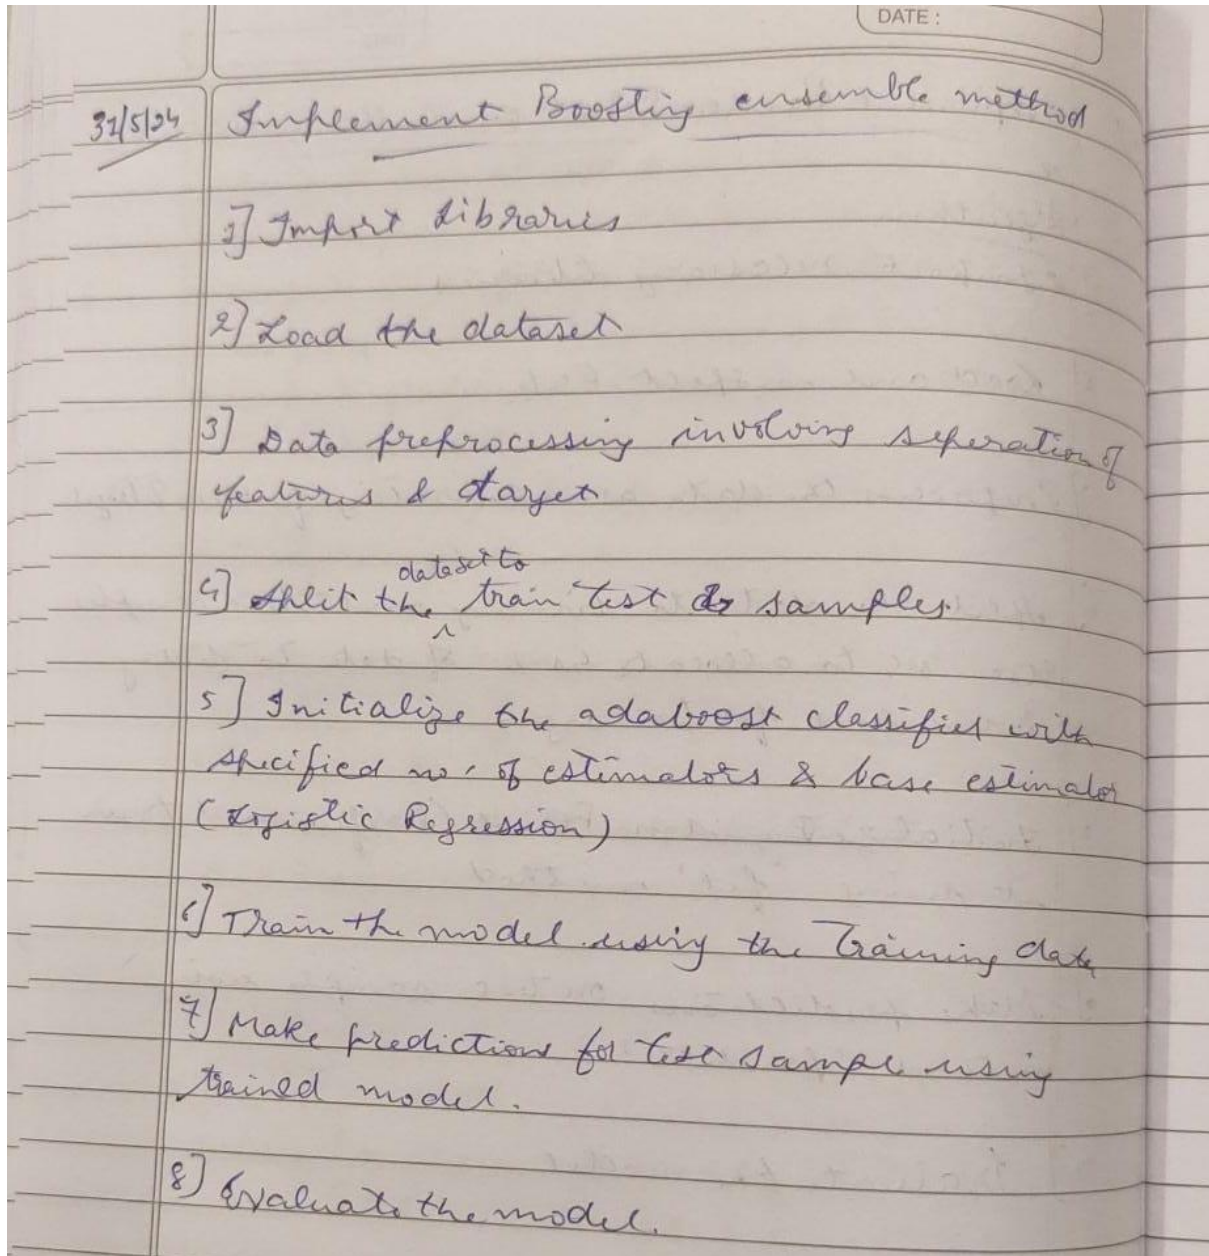
```

Confusion Matrix:
[[23  0  0]
 [ 0 19  0]
 [ 0  1 17]]

```

12) Implement Boosting ensemble method on a given dataset.

Algorithm



Code:

```

In [4]: from sklearn.linear_model import LogisticRegression
        from sklearn.ensemble import AdaBoostClassifier
        from sklearn import metrics
        from sklearn import datasets

In [5]: import pandas as pd
        import matplotlib.pyplot as plt
        from sklearn.model_selection import train_test_split

In [6]: # Load the iris dataset
        iris = datasets.load_iris()
        X = pd.DataFrame(iris.data, columns=['Sepal_Length', 'Sepal_Width', 'Petal_Length', 'Petal_Width'])
        y = pd.DataFrame(iris.target, columns=['Targets'])

In [7]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.4,random_state=42)

In [9]: mylogregmodel = LogisticRegression()

In [10]: adabc = AdaBoostClassifier(n_estimators = 150, base_estimator = mylogregmodel, learning_rate = 1)

In [11]: model = adabc.fit(X_train, y_train)

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166: FutureWarning: `base_estimator` was renamed to `estimator` in version 1.2 and will be removed in 1.4.
  warnings.warn(

In [12]: y_pred = model.predict(X_test)

In [13]: metrics.accuracy_score(y_test, y_pred)

```

Results:

```
Out[13]: 0.9833333333333333
```